Multivariate Optimization 3

Exercise 1: Stochastic Gradient Descent

Consider the ordinary linear regression (without intercept) where we want to minimize

$$\mathbb{E}_{\mathbf{x},y}\|\boldsymbol{\theta}^\top\mathbf{x} - y\|_2^2$$

with $\mathbf{x} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma_x})$ and $y \sim \mathcal{N}(\boldsymbol{\theta^*}^\top\mathbf{x}, \sigma^2)$.

(a) Show that $\mathbb{E}_{\mathbf{x},y}\nabla_{\boldsymbol{\theta}}\|\boldsymbol{\theta}^\top\mathbf{x} - y\|_2^2 = \nabla_{\boldsymbol{\theta}}\mathbb{E}_{\mathbf{x},y}\|\boldsymbol{\theta}^\top\mathbf{x} - y\|_2^2$

(b) Interpret a) in terms of SGD.

(c) Consider the the univariate setting with $\boldsymbol{\Sigma_x} = (0.5^2), \sigma = 0.1, \boldsymbol{\theta^*} = (0.5)$.
Write an R script which plots the "confusion", i.e., the variance of the gradients, for $\theta \in \{0, 0.05, 0.1, \dots 0.95, 1.0\}$ (each setting repeated 200 times).
Do two such simulation studies, one where the random batches are of size 100 and one with random batch sizes of 1000.

(d) What do you observe in c) ?

(e) Write an R script which solves the setting in c) with SGD with random batch sizes of 1 and $\alpha = 0.3$. Start with $\boldsymbol{\theta} = 0$ and do 20 iterations. (Repeat 200 times). Compare with GD.