

Optimization

Conditions for optimality



Learning goals



● TODO

● TODO

DEFINITION LOCAL AND GLOBAL MINIMUM

Given $\mathcal{S} \subseteq \mathbb{R}^d$, $f : \mathcal{S} \rightarrow \mathbb{R}$:

- f has a **global minimum** in $\mathbf{x}^* \in \mathcal{S}$, if $f(\mathbf{x}^*) \leq f(\mathbf{x})$ for all $\mathbf{x} \in \mathcal{S}$
- f has a **local minimum** in \mathbf{x}^* , if $f(\mathbf{x}^*) \leq f(\mathbf{x})$ for all $\mathbf{x} \in B_\epsilon(\mathbf{x}^*)$, with $B_\epsilon(\mathbf{x}^*) := \{\mathbf{x} \in \mathcal{S} \mid \|\mathbf{x} - \mathbf{x}^*\| < \epsilon\}$ denoting all points that have a maximum distance of ϵ to \mathbf{x}^* (“ ϵ -neighborhood”).




Source (left): https://en.wikipedia.org/wiki/Maxima_and_minima.

Source (right): <https://wngaw.github.io/linear-regression/>.

CONVEX OPTIMIZATION PROBLEMS

Why do convex optimization problems play a big role in optimization?

- In a convex optimization problem, every local optimum is a global one.
- If the objective function is strictly convex, then the problem has at most one optimal point.



Left: Example for a function that is strictly convex. It has one local minimum. Middle: A function that is convex, but not strictly convex. All local optima are global ones, but the optimum is not unique. Right: A function that is not convex.

EXISTENCE OF OPTIMA

$$f : \mathcal{S} \subseteq \mathbb{R}^d \rightarrow \mathbb{R}$$


- f continuous:
 - A real-valued function f , which is defined on a **compact set**, must attain a minimum and a maximum according to the Extreme Value Theorem.
- f not continuous:
 - In general no statement possible about existence of maximum/minimum, except on continuous intervals of the function.

Conditions for Optimality

FIRST ORDER CONDITION

There is one specific property that holds for both the univariate and multivariate case at a (local) optimum: The Taylor series approximation (i.e., the tangent hyperplane) at a local minimum is perfectly flat.

This means that at every local minima the first order derivatives are 0. This condition is also called the **first-order condition**.



Examples for a univariate (left) and multivariate (right) function. Both functions are strictly convex, and have a unique local optimum, which is also the global one. The tangent (hyperplane) is perfectly flat at the optimum.

Source: Watt, 2020, Machine Learning Refined.

FIRST ORDER CONDITION

Because at every local minimum \mathbf{x}^* the first order derivative is necessarily always zero, we call this the **first-order** or **necessary** condition.

- **First-order condition (univariate):** Let $\mathbf{x}^* \in \mathbb{R}$ be a local minimum of f . Then:

$$f'(\mathbf{x}^*) = 0$$


- **First-order condition (multivariate):** Let $\mathbf{x}^* \in \mathbb{R}^d$ be a local minimum of f . Then:

$$\nabla f(\mathbf{x}^*) = (0, 0, \dots, 0)^\top$$

The points at which the first order derivative is zero are called **stationary points**.

FIRST ORDER CONDITION

However, this is only a **necessary** condition. It is not valid the other way around, i.e. not every point at which the first order derivative is 0 (stationary point) is also a local minimum.




Left: There are for points which fulfill the necessary condition, two of them are in fact local (and also global) minima, and two are local (and also global) maxima (which is not what we are looking for!). Middle: The function has one point that fulfills the necessary condition, but it is neither a local minimum nor a local maximum; such points are called “saddle points”. Right: A function that has multiple local minima and maxima.

Source: Watt, 2020, Machine Learning Refined.

SECOND ORDER CONDITION

However, the function's second order derivative helps us identifying whether a stationary point \mathbf{x}^* is a local optimum or not. In cases where the second derivative $f''(\mathbf{x}^*) > 0$ (i.e. the function is convex in the neighborhood of \mathbf{x}^*), \mathbf{x}^* is a local minimum.



Left: Around the two minima, the function is convex; around the two maxima the function is concave. Middle: The function is neither convex nor concave around the stationary point. Right: Again, concave behavior around the maxima, and convex behavior around the minima of the function.

Source: Watt, 2020, Machine Learning Refined.

SECOND ORDER CONDITION

We introduce the second order conditions.

- **Second-order condition (univariate):** A **stationary** point $x^* \in \mathcal{S} \subseteq \mathbb{R}$ fulfills

$$f''(x^*) > 0.$$

- **Second-order condition (multivariate):** A **stationary** point $\mathbf{x}^* \in \mathcal{S} \subseteq \mathbb{R}^d$ fulfills

$$\nabla^2 f(\mathbf{x}^*) \text{ is positive semi-definitie}$$

(all eigenvalues are positive). This means the curvature is positive in all directions.

If the second-order conditions are fulfilled, then the point is a **local** minimum. The second-order condition is **sufficient** to prove that a point is a local minimum.