


# Optimization in Machine Learning

## First order methods: SGD Further Details




### Learning goals

- Decreasing step size for SGD
- Stopping rules
- SGD with momentum

# SGD WITH CONSTANT STEP-SIZE

**Example:** SGD with constant step size.



Fast convergence of SGD at beginning, erratic behavior later on (variance too big).

# SGD WITH DECREASING STEP-SIZE



- **Idea:** Decrease step size to reduce magnitude of steps of erratic behavior.
- Consider trade-off:
  - if step size  $\alpha^{[t]}$  decreases slowly, the variance of  $\nabla_x g_i(\mathbf{x})$  decreases slowly too.
  - if step size decreases too fast, we reach optimum slowly.
- SGD converges to stationary point if ratio of sum of squared step-sizes over the sum of step-sizes converges to 0:

$$\frac{\sum_{t=1}^{\infty} (\alpha^{[t]})^2}{\sum_{t=1}^{\infty} \alpha^{[t]}} = 0$$

(“how much noise affects you” vs. “how far you can get”).

# SGD WITH DECREASING STEP-SIZE

- Popular solution: step size fulfilling  $\alpha^{[t]} \in \mathcal{O}(\frac{1}{t})$ .



Example continued. Step size  $\alpha^{[t]} = \frac{0.2}{t}$ .

- Often not working well in practice: Step size really small really fast.
- Alternative:  $\alpha^{[t]} \in \mathcal{O}(\sqrt{t})$

# ADVANCED STEP-SIZE CONTROL

## Why not Armijo-based step-size control?

- Backtracking line search or other rules based on the Armijo-condition are usually not suitable: A check of Armijo condition

$$f(\mathbf{x} + \alpha \mathbf{d}) \leq f(\mathbf{x}) + \gamma \alpha \nabla g(\mathbf{x})^\top \mathbf{d}$$

requires evaluating the full gradient.


- But: SGD is particularly used to **avoid** expensive evaluations of gradient!
- Recent research aims at finding approximate line-search methods that provide better convergence methods, e.g., *Vaswani et al., Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates (NeurIPS 2019)*.

# MINI-BATCHES

- Reduce noise by increasing batch size  $m$  for better approximation

$$\hat{\mathbf{d}} = \frac{1}{m} \sum_{i \in J} \nabla_{\mathbf{x}} g_i(\mathbf{x}) \approx \frac{1}{n} \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) = \mathbf{d}$$

- Usually, the batch size is limited by computational resources (e.g., how much data you can load into the memory)



Example continued. Batch size  $m = 1$  vs.  $m = 5$ .

# STOPPING RULES FOR SGD

- **For GD:** We usually stop when gradient is close to 0 (i.e., we are close to a stationary point)
- **For SGD:** individual gradients do not necessarily go to zero, and we cannot access full gradient.
- Practicable solution for ML:
  - Measure the validation set error after  $T$  iterations
  - Stop, if validation set error is not improving

# SGD AND ML

## General remarks:

- SGD is a simplification of GD
- SGD particularly suitable for large-scale ML when the evaluating the gradient is too expensive / restricted by computational resources
- SGD and variants are the most commonly used methods in modern ML, for example:

- Linear models

Note that even for the linear model and quadratic loss, where a closed form solution is available, SGD might be used if the size  $n$  of the dataset is too large and the design matrix does not fit into memory.

- Neural networks
- Support vector machines
- ...



# SGD WITH MOMENTUM

SGD is usually used with momentum due to reasons mentioned in previous chapters.

---

**Algorithm 1** Stochastic gradient descent with momentum

---

- 1: **require** learning rate  $\alpha$  and momentum  $\varphi$
  - 2: **require** initial parameter  $\mathbf{x}$  and initial velocity  $\boldsymbol{\nu}$
  - 3: **while** stopping criterion not met **do**
  - 4:     Sample a minibatch of  $m$  examples from the training set  $\{\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(m)}\}$
  - 5:     Compute gradient estimate:  $\nabla \hat{f}(\mathbf{x})$
  - 6:     Compute velocity update:  $\boldsymbol{\nu} \leftarrow \varphi \boldsymbol{\nu} - \alpha \nabla \hat{f}(\mathbf{x})$
  - 7:     Apply update:  $\mathbf{x} \leftarrow \mathbf{x} + \boldsymbol{\nu}$
  - 8: **end while**
-