


Optimization in Machine Learning

Coordinate descent



Learning goals

- Axes as descent direction
- CD on linear model and LASSO
- Soft-Thresholding operator

COORDINATE DESCENT

If derivative of objective function does not exist / is unknown we cannot compute a descent direction analytically and an **inexact** procedure must be used.

Idea: Use direction of coordinate axes as “descent directions”.

In the simplest case we run iteratively over all coordinates $\{1, \dots, d\}$ and minimize f with respect to the corresponding dimension.

COORDINATE DESCENT

- First a starting point $\mathbf{x}^{[0]} = (x_1^{[0]}, \dots, x_d^{[0]})$ is selected.
- In step t we search the value x_i for each dimension $i \in \{1, 2, \dots, d\}$ that minimizes f , given $x_1^{[t]}, \dots, x_{i-1}^{[t]}$ and $x_{i+1}^{[t-1]}, \dots, x_d^{[t-1]}$:


$$x_1^{[t]} = \arg \min_{x_1} f(x_1, x_2^{[t-1]}, x_3^{[t-1]}, \dots, x_d^{[t-1]})$$

$$x_2^{[t]} = \arg \min_{x_2} f(x_1^{[t]}, x_2, x_3^{[t-1]}, \dots, x_d^{[t-1]})$$

$$x_3^{[t]} = \arg \min_{x_3} f(x_1^{[t]}, x_2^{[t]}, x_3, \dots, x_d^{[t-1]})$$

\vdots

$$x_d^{[t]} = \arg \min_{x_d} f(x_1^{[t]}, x_2^{[t]}, x_3^{[t]}, \dots, x_d)$$



https://commons.wikimedia.org/wiki/File:Coordinate_descent.svg

COORDINATE DESCENT

- Minimum is determined with (exact / inexact) line search
- Order in which the dimensions are gone through can be any permutation of $\{1, 2, \dots, d\}$
- **Convergence:** if $f(\cdot)$ is continuously differentiable and the univariate minimization problems have unique solutions, the sequence $\mathbf{x}^{[t]}$ converges to \mathbf{x}^* with $\nabla f(\mathbf{x}^*) = 0$.

The following holds:

$$f(\mathbf{x}^{[0]}) \geq f(\mathbf{x}^{[1]}) \geq f(\mathbf{x}^{[2]}) \geq \dots$$

EXAMPLE: LINEAR REGRESSION

Minimize LM with L2-loss via CD:

$$\min_{\theta} g(\theta) = \min_{\theta} \frac{1}{2} \sum_{i=1}^n \left(y^{(i)} - \theta^{\top} \mathbf{x}^{(i)} \right)^2 = \min_{\theta} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\theta\|^2$$

where $\mathbf{y} \in \mathbb{R}^n$, $\mathbf{X} \in \mathbb{R}^{n \times p}$ with columns $\mathbf{X}_1, \dots, \mathbf{X}_p \in \mathbb{R}^n$.

Assumption: data is scaled $\mathbf{X}_j^{\top} \mathbf{1} = 0$ and $\mathbf{X}^{\top} \mathbf{X} = \mathbf{1}_p$.

g simplifies to

$$\begin{aligned} g(\theta) &= \frac{1}{2} \mathbf{y}^{\top} \mathbf{y} + \frac{1}{2} \theta^{\top} \theta - \mathbf{y}^{\top} \mathbf{X} \theta \\ &\stackrel{(*)}{=} \frac{1}{2} \mathbf{y}^{\top} \mathbf{y} + \frac{1}{2} \theta^{\top} \theta - \sum_{k=1}^p \mathbf{y}^{\top} \mathbf{X}_k \theta_k \end{aligned}$$

$$(*) \quad \mathbf{X} \theta = \mathbf{X}_1 \theta_1 + \mathbf{X}_2 \theta_2 + \dots + \mathbf{X}_p \theta_p = \sum_{k=1}^p \mathbf{X}_k \theta_k.$$

EXAMPLE: LINEAR REGRESSION

To compute the exact CD update in direction j we compute

$$\frac{\partial g(\theta)}{\partial \theta_j} = \theta_j - \mathbf{y}^\top \mathbf{x}_j$$

By solving $\frac{\partial g(\theta)}{\partial \theta_j} = 0$, we get

$$\theta_j^* = \mathbf{y}^\top \mathbf{x}_j$$

as exact update for CD in direction j . We repeat this update over all variables.

SOFT THRESHOLDING OPERATOR

Minimize LM with L2-loss and L1 regularization via CD:

$$\min_{\boldsymbol{\theta}} h(\boldsymbol{\theta}) = \min_{\boldsymbol{\theta}} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|^2 + \lambda \|\boldsymbol{\theta}\|_1$$

We can write $h(\boldsymbol{\theta}) = \frac{1}{2} \mathbf{y}^\top \mathbf{y} + \frac{1}{2} \boldsymbol{\theta}^\top \boldsymbol{\theta} - \sum_{k=1}^p (\mathbf{y}^\top \mathbf{X}_k \boldsymbol{\theta}_k + \lambda |\boldsymbol{\theta}_k|)$.

Because $|\cdot|$ is not differentiable, we distinguish three cases:

- **Case 1:** $\boldsymbol{\theta}_j > 0$. Then $|\boldsymbol{\theta}_j| = \boldsymbol{\theta}_j$ and

$$0 = \frac{\partial g(\boldsymbol{\theta})}{\partial \boldsymbol{\theta}_j} = \boldsymbol{\theta}_j - \mathbf{y}^\top \mathbf{X}_j + \lambda \quad \Leftrightarrow \quad \boldsymbol{\theta}_{j,\text{LASSO}}^* = \boldsymbol{\theta}_j^* - \lambda$$

- **Case 2:** $\boldsymbol{\theta}_j < 0$. Then $|\boldsymbol{\theta}_j| = -\boldsymbol{\theta}_j$ and

$$0 = \frac{\partial g(\boldsymbol{\theta})}{\partial \boldsymbol{\theta}_j} = \boldsymbol{\theta}_j - \mathbf{y}^\top \mathbf{X}_j - \lambda \quad \Leftrightarrow \quad \boldsymbol{\theta}_{j,\text{LASSO}}^* = \boldsymbol{\theta}_j^* + \lambda$$

- **Case 3:** $\boldsymbol{\theta}_j = 0$.

SOFT THRESHOLDING OPERATOR

We can write the solution as:

$$\theta_{j,\text{LASSO}}^* = \begin{cases} \theta_j^* - \lambda & \text{if } \theta_j^* > \lambda \\ \theta_j^* + \lambda & \text{if } \theta_j^* < -\lambda \\ 0 & \text{if } \theta_j^* \in [-\lambda, \lambda], \end{cases}$$

which is also referred to as **soft-thresholding operator**. Coefficients for which the solution to the unregularized problem are smaller than a threshold, $|\theta_j^*| < \lambda$, are shrunk to zero.

Note:

- For case 1, we require

$$\theta_{j,\text{LASSO}}^* = \theta_j^* - \lambda > 0 \quad \Leftrightarrow \quad \theta_j^* > \lambda$$

- For case 2, we require

$$\theta_{j,\text{LASSO}}^* = \theta_j^* + \lambda < 0 \quad \Leftrightarrow \quad \theta_j^* < -\lambda$$

CD FOR STATISTICS AND ML

Why is it being used?

- Very easy to implement.
- Good implementation can achieve state-of-the-art performance.
- Scalable, e.g. no storage or operations on large objects, only the current point
- Applicable in both differentiable and derivative-free cases.

Examples:

- Lasso regression, Lasso GLM, graphical Lasso
- Support Vector Machines
- Regression with non-convex penalties